

DDIA



Project: Weather Stations Monitoring.



Name	ID
Karim Fathy	20011116
Mohamed Amr	20011675
Omar Mahmoud	20011027
Omar Tarek	20010998

Project main points:

Bitcask

Overview

This Java-based Bitcask server is a lightweight, concurrent TCP server designed to handle read (r), write (w), and append (a) operations from clients. It supports multi-threaded request processing and asynchronous log tracking for durability and performance optimization.

Key Components

1. Server Initialization

- The server reads configuration from a `system.properties` file, extracting:
 - `server.port`: TCP port to listen on.
 - `server.logs`: Path to the log file(s).
 - `server.hints`: Path to hint files (for indexing/optimization).
 - `server.threads`: Number of threads used to process requests.

2. Client Handling (`handleClient` method)

- Each client connection is handled in a separate thread.
- Uses a fixed thread pool (`ExecutorService`) to manage and execute Worker tasks.
- Accepts commands prefixed with:
 - w: Write
 - r: Read
 - a: Read All
- Supports the `--no-reply` flag to execute commands without awaiting a response.
- Valid requests are submitted to a Worker class for processing.
- Results (for requests expecting replies) are collected and sent back to the client.

3. Worker Execution

- Each Worker instance represents a task for handling a single client command.
- Asynchronous execution allows high throughput and non-blocking processing.

4. Log Tracking

- LogPathTracker is initialized to monitor the log directory.
- It polls at a fixed interval (every 5 seconds) and likely handles log file rotation or cleanup.
- Helps in managing storage and maintaining performance.

5. Server Loop

- The main thread continuously listens for incoming TCP connections.
- Upon accepting a new connection, it delegates client handling to a new thread using `handleClient`.

Concurrency & Fault Tolerance

- Utilizes Java's `ExecutorService` for concurrent request execution.
- Separate threads for each client ensure isolation and responsiveness.
- Graceful shutdown of executors and sockets prevents resource leaks.
- Exceptions are logged, and processing continues robustly.

Client Functionality (BitcaskClient)

- **Core Role:** Provides CLI-based interaction with the Bitcask server.
 - **Commands Supported:**
 - `--view-all`: Retrieves all stored data as a CSV file.
 - `--view --key=SOME_KEY`: Retrieves a value for a specific key.
 - `--perf --clients=N`: Launches N concurrent client threads to perform `--view-all` operations for performance benchmarking.
 - **Implementation Highlights:**
 - Uses Socket programming to connect to the server on `localhost:5000` (modifiable).
 - Exports `--view-all` results to timestamped CSV files at a specified file path.
 - `--perf` tests simulate concurrent clients to evaluate server throughput.
-

KafkaToBitcask (Normalizer)

Overview

This component acts as a **bridge between Kafka and the Bitcask server**, consuming streaming data from a Kafka topic (weather-station), transforming it as needed, and sending each key-value record to the Bitcask server over TCP.

Component Role

- **Input:** Kafka topic weather-station (stream of key-value string pairs).
- **Output:** TCP-based w (write) command to the Bitcask server.
- **Purpose:** Normalize incoming Kafka messages and persist them in a Bitcask-inspired data store.

Configuration

- **Kafka Streams App ID:** bitcask-producer
- **Kafka Bootstrap Servers:** localhost:9092
- **SerDes:** String SerDes used for both key and value.

Stream Processing Logic

```
KStream<String, String> stream = builder.stream("input-topic");
stream.foreach((key, value) -> {
    Long id = Long.parseLong(key);
    sendToBitcask(id, value, "localhost", 5000);
});
```

- **Key Parsing:** Converts key from String to Long.
 - **Value Normalization:** Replaces commas in values with semicolons to avoid CSV conflicts.
 - **Command Sent:** w <key> <value>--no-reply
-

WeatherStatusIndexer (Elasticsearch)

Overview

This Python class reads weather data in .parquet format from a directory of station folders, flattens nested fields, tags dropped messages, and indexes the results into an **Elasticsearch** index called "weather_statuses".

Component Role

- **Input:** Parquet files under root_data_dir/stationId_x/ folders.
- **Output:** Documents indexed in Elasticsearch.
- **Goal:** Store real-time or batch IoT weather data for analysis and querying.

Main Functionalities

1. flatten_and_tag(df, station_id)

- **Flattens** the nested "weather" field.
- **Adds:**
 - station_id
 - status_timestamp as datetime
 - dropped flag: True if sequence number (s_no) skips.
- **Sorts** by s_no to detect dropped messages.

2. index_exists_or_create()

- Deletes the index if it already exists.
- Creates a new index with mappings:
 - station_id, s_no: long
 - battery_status: keyword
 - status_timestamp: date
 - humidity, temperature, wind_speed: integer
 - dropped: boolean

3. process_all_stations()

- Iterates through all subfolders.
- Extracts station ID from folder name like stationId_3.

- Reads all .parquet files.
- Applies flattening and tagging.
- Aggregates all documents.

4. bulk_upload(docs)


- Uses helpers.bulk() from elasticsearch-py to send all docs.

5. run()

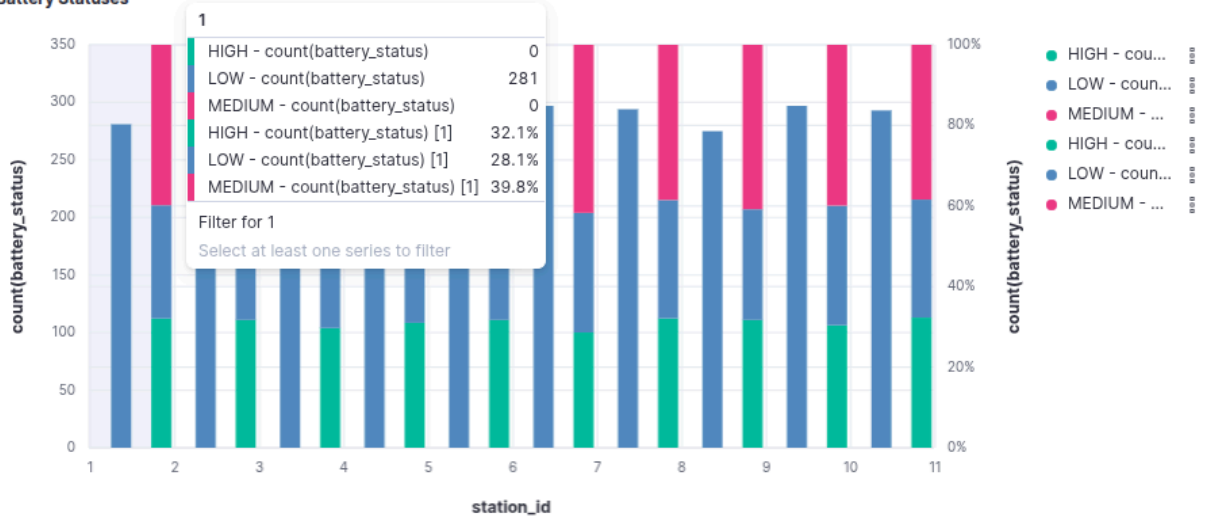
- Calls the methods in sequence: index creation → processing → upload.

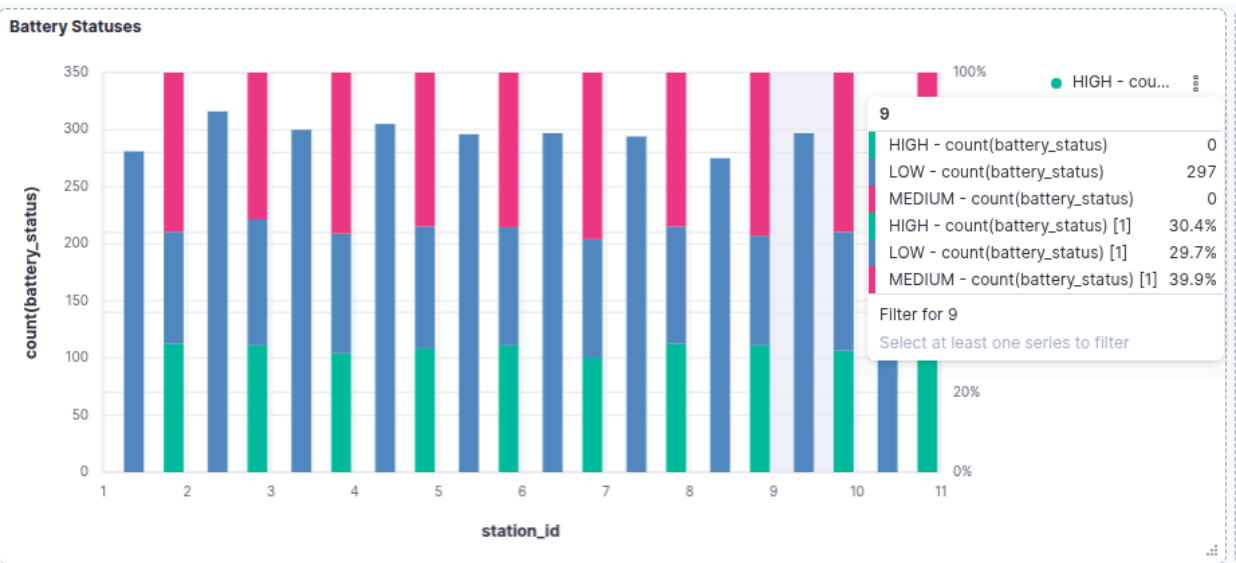
Kibana visualization

Battery Statuses




Battery Statuses



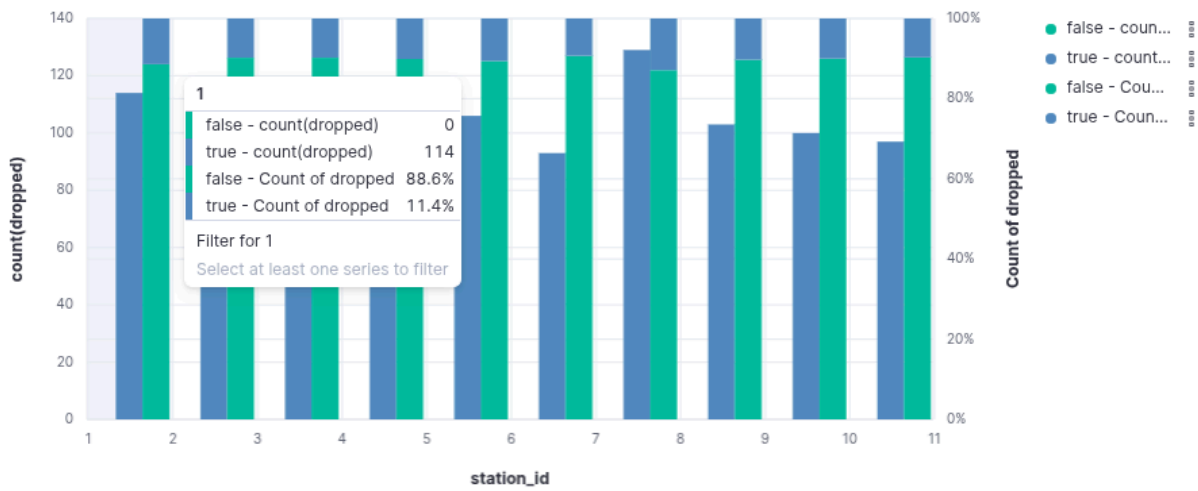



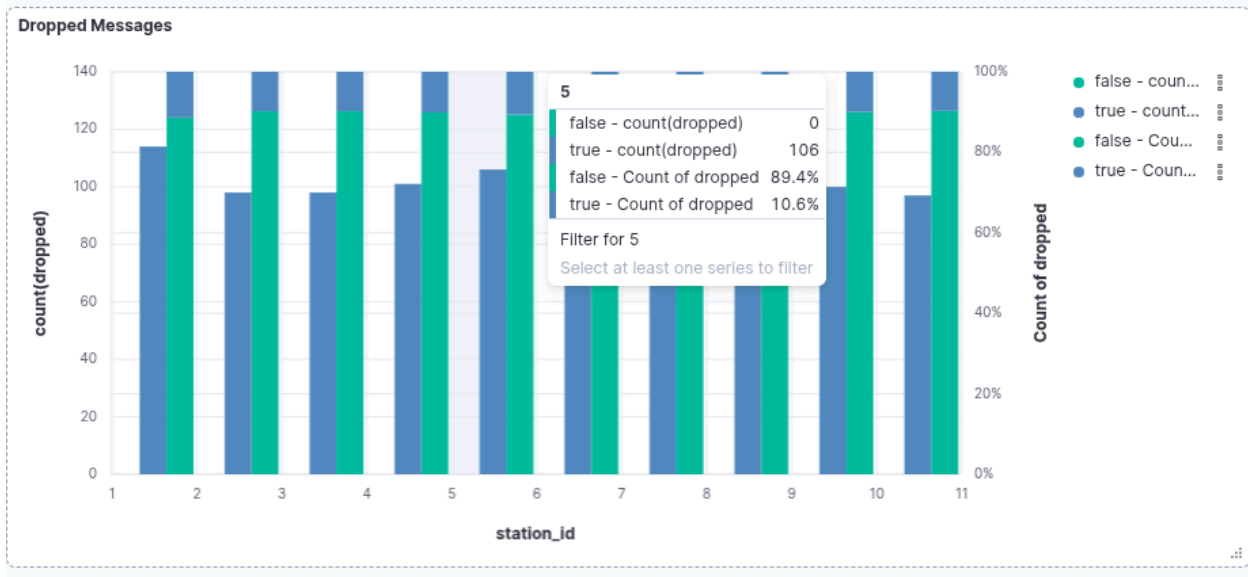
Dropped Messages

Dropped Messages



Dropped Messages








BitCask Riak LSM directory

Structure

```
|
|--> logs —> 0, 1, 2, 3, 4
|--> hints —> 0, 1, 2, 3, 4
```



Example parquet file:

s_no ?1	battery_status ?1	status_timestamp ?1	weather ?1
1110	HIGH	1748219344	{"humidity":50,"temperature":65,"win...
1111	MEDIUM	1748219345	{"humidity":50,"temperature":65,"win...
1113	MEDIUM	1748219346	{"humidity":50,"temperature":65,"win...
1114	HIGH	1748219347	{"humidity":50,"temperature":65,"win...
1115	HIGH	1748219348	{"humidity":50,"temperature":65,"win...
1116	LOW	1748219350	{"humidity":50,"temperature":65,"win...
1117	LOW	1748219351	{"humidity":50,"temperature":65,"win...
1118	HIGH	1748219352	{"humidity":50,"temperature":65,"win...
1122	MEDIUM	1748219353	{"humidity":50,"temperature":65,"win...
1123	MEDIUM	1748219354	{"humidity":50,"temperature":65,"win...
1125	HIGH	1748219355	{"humidity":50,"temperature":65,"win...
1126	HIGH	1748219356	{"humidity":50,"temperature":65,"win...
1127	MEDIUM	1748219357	{"humidity":50,"temperature":65,"win...
1128	HIGH	1748219358	{"humidity":50,"temperature":65,"win...
1129	HIGH	1748219359	{"humidity":50,"temperature":65,"win...
1130	MEDIUM	1748219360	{"humidity":50,"temperature":65,"win...
1131	HIGH	1748219361	{"humidity":50,"temperature":65,"win...
1133	LOW	1748219362	{"humidity":50,"temperature":65,"win...
1134	MEDIUM	1748219363	{"humidity":50,"temperature":65,"win...
scottpaulin ✕	HIGH	1748219364	{"humidity":50,"temperature":65,"win...

Used online viewer to view parquet file in columns

Source Code:

Weather station:

```
public class Main {
    private static final Random rand = new Random();

    public static void main(String[] args) throws InterruptedException {

        long stationId = 1;
        String stID = System.getenv("STATION_ID");
        // Check if station ID is in the format "weather-station-X"
        if (stID != null && stID.matches("weather-station-\\d+")) {
            try {
                String idStr = stID.substring(stID.lastIndexOf('-') + 1);
                stationId = Long.parseLong(idStr)+1;
                System.out.println("Found weather-station format ID: " + stationId);
            } catch (NumberFormatException e) {
                System.err.println("Failed to parse station number from: " + stID);
            }
        }
        System.out.println("Station ID: " + stationId);

        long statusMsgCounter = 0;

        String bootstrapServers = System.getenv("KAFKA_BOOTSTRAP_SERVERS");
        if (bootstrapServers == null || bootstrapServers.isEmpty()) {
            bootstrapServers = "localhost:9092"; // Default to Kubernetes service name
        }

        Properties props = new Properties();
        props.setProperty(ProducerConfig.BootstrapServersConfig, bootstrapServers);
        props.setProperty(ProducerConfig.KeySerializerClassConfig, StringSerializer.class.getName());
        props.setProperty(ProducerConfig.ValueSerializerClassConfig,
StringSerializer.class.getName());
        System.out.println("Connecting to Kafka at: " + bootstrapServers);

        ObjectMapper objectMapper = new ObjectMapper();
        try(Producer<String,String> producer = new KafkaProducer<>(props)){
            System.out.println("Weather station " + stationId + " started. Sending data to Kafka...");

            while(true){
                if (dropMsg()){
                    statusMsgCounter++;
                }
            }
        }
    }
}
```

```

        System.out.println("Message dropped");
        continue;
    }
    WeatherStationMsg msg = createWeatherStationMsg(stationId, statusMsgCounter++);
    try {
        String jsonMsg = objectMapper.writeValueAsString(msg);
        ProducerRecord<String,String> record = new ProducerRecord<>("weather-station",
String.valueOf(stationId), jsonMsg);
        producer.send(record, (metadata, exception) -> {
            if (exception != null) {
                System.out.println("Error sending message: " + exception.getMessage());
            } else {
                System.out.println("Message sent to topic " + metadata.topic() + " partition
" + metadata.partition() + " offset " + metadata.offset());
            }
        });
    } catch (JsonProcessingException e) {
        System.out.println("message to json error");
        throw new RuntimeException(e);
    }
    Thread.sleep(1000);
}
}

private static boolean dropMsg(){
    return rand.nextInt(100) < 10; // 10% chance to drop the message
}

private static WeatherStationMsg createWeatherStationMsg(long stationId, long statusMsgCounter) {

    int temperature = rand.nextInt(40,100) ;
    int humidity = rand.nextInt(101);
    int windSpeed = rand.nextInt(26);

    WeatherData weatherData = new WeatherData(temperature, humidity, windSpeed);
    int batteryLevel = rand.nextInt(101);

    BatteryStatus batteryStatus ;
    if (batteryLevel>=70){
        batteryStatus = BatteryStatus.HIGH;
    } else if (batteryLevel>=30){
        batteryStatus = BatteryStatus.MEDIUM;
    } else {
        batteryStatus = BatteryStatus.LOW;
    }
    return new WeatherStationMsg(stationId, statusMsgCounter, batteryStatus, weatherData);
}

```

```
}  
}
```

Model classes for message creation and serialization to json:

```
package org.example;  
  
import java.time.Instant;  
  
public class WeatherStationMsg {  
    private long station_id;  
    private long s_no;//auto increment  
    private BatteryStatus battery_status;  
    private long status_timestamp;  
    private WeatherData weather;  
  
    public WeatherStationMsg(long stationId,long s_no,BatteryStatus batteryStatus, WeatherData  
weatherData) {  
        this.station_id = stationId;  
        this.s_no = s_no;  
        this.battery_status = batteryStatus;  
        this.status_timestamp = Instant.now().getEpochSecond();  
        this.weather = weatherData;  
    }  
  
    public long getStation_id() {  
        return station_id;  
    }  
  
    public void setStation_id(long station_id) {  
        this.station_id = station_id;  
    }  
  
    public long getS_no() {  
        return s_no;  
    }  
  
    public void setS_no(long s_no) {  
        this.s_no = s_no;  
    }  
  
    public BatteryStatus getBattery_status() {  
        return battery_status;  
    }  
}
```

```

}

public void setBattery_status(BatteryStatus battery_status) {
    this.battery_status = battery_status;
}

public long getStatus_timestamp() {
    return status_timestamp;
}

public void setStatus_timestamp(long status_timestamp) {
    this.status_timestamp = status_timestamp;
}

public WeatherData getWeather() {
    return weather;
}

public void setWeather(WeatherData weather) {
    this.weather = weather;
}
}

```

```

package org.example;

public class WeatherData {
    private final int humidity;
    private final int temperature;
    private final int wind_speed;

    public WeatherData(int temperature, int humidity, int windSpeed) {
        this.temperature = temperature;
        this.humidity = humidity;
        this.wind_speed = windSpeed;
    }

    public int getHumidity() {
        return humidity;
    }

    public int getTemperature() {
        return temperature;
    }
}

```



```
    public int getWind_speed() {  
        return wind_speed;  
    }  
}
```

```
package org.example;
```

```
public enum BatteryStatus {  
    HIGH,  
    MEDIUM,  
    LOW  
}
```

Dockerfile to create weather station image:

```
FROM eclipse-temurin:21-jdk-alpine

WORKDIR /app

# Copy the pre-built JAR file from your local filesystem
COPY target/weather_station-1.0-SNAPSHOT-jar-with-dependencies.jar app.jar

# Default environment variables
ENV KAFKA_BOOTSTRAP_SERVERS=kafka:9092
ENV STATION_ID=1

CMD ["java", "-jar", "app.jar"]
```

Central station:

1) Kafka message to bitcask

```
public class KafkaToBitcask {

    public static void main(String[] args) {

        String KAFKA_BOOTSTRAP_SERVERS = System.getenv("KAFKA_BOOTSTRAP_SERVERS");
        if (KAFKA_BOOTSTRAP_SERVERS == null) {
            System.err.println("Environment variable KAFKA_BOOTSTRAP_SERVERS is not set.");
            System.exit(1);
        }
        String KAFKA_TOPIC = System.getenv("KAFKA_TOPIC");
        if (KAFKA_TOPIC == null) {
            System.err.println("Environment variable KAFKA_TOPIC is not set.");
            System.exit(1);
        }

        String BITCASK_SERVER_HOST = System.getenv("BITCASK_SERVER_HOST");
        if (BITCASK_SERVER_HOST == null) {
            System.err.println("Environment variable BITCASK_SERVER_HOST is not set.");
            System.exit(1);
        }

        String BITCASK_SERVER_PORT_str = System.getenv("BITCASK_SERVER_PORT");
        if (BITCASK_SERVER_PORT_str == null) {
            System.err.println("Environment variable BITCASK_SERVER_PORT is not set.");
            System.exit(1);
        }
        int BITCASK_SERVER_PORT = Integer.parseInt(BITCASK_SERVER_PORT_str);
        // Check if the environment variables are set
        if (KAFKA_BOOTSTRAP_SERVERS.isEmpty() || BITCASK_SERVER_HOST.isEmpty() ||
        BITCASK_SERVER_PORT_str.isEmpty()) {
            System.err.println("Environment variables are not set.");
            System.exit(1);
        }
        // Kafka Streams Configuration
        Properties props = new Properties();
        props.put(StreamsConfig.APPLICATION_ID_CONFIG, "message-normalizer");
        props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, KAFKA_BOOTSTRAP_SERVERS);
        props.put(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass());
        props.put(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG, Serdes.String().getClass());
```

```

// Define Stream Processing
StreamsBuilder builder = new StreamsBuilder();
KStream<String, String> stream = builder.stream(KAFKA_TOPIC);

stream.foreach((key, value) -> {
    try {
        Long id = Long.parseLong(key);

        // Send to Bitcask server via TCP
        sendToBitcask(id, value, BITCASK_SERVER_HOST, BITCASK_SERVER_PORT); // change host/port
as needed
    } catch (Exception e) {
        e.printStackTrace();
    }
});

KafkaStreams streams = new KafkaStreams(builder.build(), props);
streams.start();

Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
}

public static void sendToBitcask(Long key, String value, String host, int port) throws IOException {
    try (Socket socket = new Socket(host, port);
        PrintWriter writer = new PrintWriter(
            new OutputStreamWriter(socket.getOutputStream(), StandardCharsets.UTF_8), true)){

        value = value.replace(',', ';');
        String request = "w" + " " + key + " " + value + "--no-reply";
        writer.println(request);

    } catch (IOException e) {
        e.printStackTrace();
    }
}
}

```

2) Rain-detector:

```
public class Main {
    public static void main(String[] args) {

        String KAFKA_BOOTSTRAP_SERVERS = System.getenv("KAFKA_BOOTSTRAP_SERVERS");
        if (KAFKA_BOOTSTRAP_SERVERS == null) {
            KAFKA_BOOTSTRAP_SERVERS = "localhost:9092";
        }

        Properties props = new Properties();
        props.setProperty(StreamsConfig.APPLICATION_ID_CONFIG, "rain-detector");
        props.setProperty(StreamsConfig.BootstrapServersConfig, KAFKA_BOOTSTRAP_SERVERS);

        props.setProperty(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass().getName());
        props.setProperty(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG, Serdes.String().getClass().getName());

        ObjectMapper objectMapper = new ObjectMapper();

        StreamsBuilder builder = new StreamsBuilder();

        KStream<String, String> inputStream = builder.stream("weather-station");

        KStream<String, String> rainStream = inputStream.filter((key, value) -> {
            try {
                JsonNode jsonNode = objectMapper.readTree(value);
                int humidity = jsonNode.get("weather").get("humidity").asInt();
                return humidity >= 70;
            } catch (Exception e) {
                e.printStackTrace();
                return false;
            }
        }).mapValues((value) ->{
            try{

                JsonNode jsNode = objectMapper.readTree(value);
                int stationId = jsNode.get("station_id").asInt();
                int humidity = jsNode.get("weather").get("humidity").asInt();
                String s = "Rain detected at station ID: " + stationId + " at humidity level: " +
humidity;

                System.out.println(s);
                return s;
            }
        });
    }
}
```

```
    }  
    catch(Exception e){  
        e.printStackTrace();  
        return "Error processing message";  
    }  
});  
  
rainStream.to("rain-alerts");  
  
KafkaStreams streams = new KafkaStreams(builder.build(), props);  
streams.start();  
Runtime.getRuntime().addShutdownHook(new Thread(streams::close));  
}  
}
```

3) Parquet-maker:

```
public class KafkaToParquetWriter {

    private final String bootstrapServers;
    private final String topic;
    private final String groupId;
    private final String outputDir;
    private final long maxFileSizeBytes;
    private final int batchSize;
    private final Schema schema;
    private long estimatedRecordSize;
    private final AtomicBoolean running = new AtomicBoolean(true);

    // Keep track of active writers and their current sizes
    private final Map<String, WriterInfo> activeWriters = new ConcurrentHashMap<>();
    private final Map<String, Long> lastActivityTimestamp = new ConcurrentHashMap<>();
    private final Map<String, List<GenericRecord>> recordBatches = new ConcurrentHashMap<>();
    private final long inactivityThresholdMs;

    // Track file sequence numbers for each partition
    private final Map<String, Integer> partitionFileCounters = new ConcurrentHashMap<>();

    public KafkaToParquetWriter(String bootstrapServers, String topic, String groupId,
                               String outputDir, Schema schema, long maxFileSizeBytesMB, int batchSize) {
        this.bootstrapServers = bootstrapServers;
        this.topic = topic;
        this.groupId = groupId;
        this.outputDir = outputDir;
        this.schema = schema;
        this.maxFileSizeBytes = maxFileSizeBytesMB * 1024 * 1024; // Not used with batch-per-file
        approach
        this.batchSize = batchSize; // Each batch gets its own file
        this.estimatedRecordSize = -1; // Initialize estimated record size
        this.inactivityThresholdMs = 300000; // 5 minutes inactivity threshold

        // Create output directory if it doesn't exist
        File directory = new File(outputDir);
        if (!directory.exists()) {
            directory.mkdirs();
        }
    }

    public void start() {
        // Create Kafka consumer
```

```

Properties props = new Properties();
props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, bootstrapServers);
props.put(ConsumerConfig.GROUP_ID_CONFIG, groupId);
props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, StringDeserializer.class.getName());
props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, StringDeserializer.class.getName());
props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

// Start a timer to periodically check for inactive partitions
Timer inactivityTimer = new Timer("InactivityChecker", true);
inactivityTimer.scheduleAtFixedRate(new TimerTask() {
    @Override
    public void run() {
        checkInactivePartitions();
    }
}, 300000, 300000); // Check every 5 min

try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props)) {
    consumer.subscribe(Collections.singletonList(topic));

    // Main polling loop
    // Main polling loop
    while (running.get()) {
        ConsumerRecords<String, String> records = consumer.poll(Duration.ofMillis(100));
        boolean newMsgs = false;

        for (ConsumerRecord<String, String> record : records) {
            try {
                newMsgs = true;
                // Parse the JSON message
                JSONObject json = new JSONObject(record.value());

                // Extract timestamp and station ID
                String stationIdStr = json.optString("station_id");

                // Handle timestamp - could be either Unix timestamp (long) or ISO formatted
                string
                LocalDateTime dateTime;
                if (json.has("status_timestamp")) {
                    Object timestampObj = json.get("status_timestamp");
                    if (timestampObj instanceof Long || timestampObj instanceof Integer) {
                        // Handle numeric timestamp (Unix timestamp in seconds)
                        long timestamp = json.getLong("status_timestamp");
                        dateTime = LocalDateTime.ofEpochSecond(timestamp, 0,
java.time.ZoneOffset.UTC);
                    } else {

```



```

        // Handle string timestamp in ISO format
        String timestamp = json.getString("status_timestamp");
        try {
            dateTime = LocalDateTime.parse(timestamp,
DateTimeFormatter.ISO_DATE_TIME);
        } catch (Exception e) {
            // Fallback - use current time if parsing fails
            System.err.println("Could not parse timestamp: " + timestamp + ",
using current time");

            dateTime = LocalDateTime.now();
        }
    } else {
        // No timestamp in the message, use current time
        System.err.println("Message missing timestamp, using current time");
        dateTime = LocalDateTime.now();
    }

    if (stationIdStr.isEmpty()) {
        System.err.println("Message missing station_id field: " + record.value());
        continue;
    }

    String datePath = dateTime.format(DateTimeFormatter.ofPattern("yyyy/MM/dd"));

    // Swap the partitioning order: stationId/datePath instead of datePath/stationId
    String partitionKey = stationIdStr + "/" + datePath.replace("/", "_");

    // Update last activity timestamp for this partition
    lastActivityTimestamp.put(partitionKey, System.currentTimeMillis());

    // Add to appropriate batch
    recordBatches.computeIfAbsent(partitionKey, k -> new ArrayList<>())
        .add(convertJsonToAvro(json));

    // Process batch if it reaches the batch size
    List<GenericRecord> batch = recordBatches.get(partitionKey);
    if (batch.size() >= batchSize) {
        writeToParquet(partitionKey, batch);
        recordBatches.put(partitionKey, new ArrayList<>());
    }
} catch (Exception e) {
    System.err.println("Error processing record: " + e.getMessage());
    e.printStackTrace();
}

```

```

        }
        if (newMsgs){
            consumer.commitSync();
        }
    }

    // Clean up timer when shutting down
    inactivityTimer.cancel();

    // Close all active writers when shutting down
    closeAllWriters();
} catch (Exception e) {
    System.err.println("Error in Kafka consumer: " + e.getMessage());
    e.printStackTrace();
}
}

private GenericRecord convertJsonToAvro(JSONObject json) {
    GenericRecord avroRecord = new GenericData.Record(schema);

    // Map JSON fields to Avro schema fields - handle nested structures
    for (Schema.Field field : schema.getFields()) {
        String fieldName = field.name();
        if (json.has(fieldName)) {
            if (fieldName.equals("weather") && !json.isNull("weather")) {
                // Handle nested weather object
                JSONObject weatherJson = json.getJSONObject("weather");
                GenericRecord weatherRecord = new GenericData.Record(field.schema());

                // Map weather fields
                for (Schema.Field weatherField : field.schema().getFields()) {
                    String weatherFieldName = weatherField.name();
                    if (weatherJson.has(weatherFieldName)) {
                        weatherRecord.put(weatherFieldName, weatherJson.get(weatherFieldName));
                    }
                }

                avroRecord.put(fieldName, weatherRecord);
            } else {
                // Handle regular fields
                avroRecord.put(fieldName, json.get(fieldName));
            }
        }
    }
}

```

```
    return avroRecord;
```

```
}
```

```
private synchronized void writeToParquet(String partitionKey, List<GenericRecord> records) throws  
IOException {
```

```
    if (records.isEmpty()) {
```

```
        return;
```

```
    }
```

```
    // Get and increment the file counter for this partition
```

```
    int fileId = partitionFileCounters.compute(partitionKey, (key, count) -> count == null ? 1 :  
count + 1);
```

```
    // Create a new file name for this batch
```

```
    String fileName = "part-" + String.format("%06d", fileId) + ".parquet";
```

```
    String filePath = outputDir + "/" + partitionKey + "/" + fileName;
```

```
    // Ensure directory exists
```

```
    File directory = new File(new File(filePath).getParent());
```

```
    if (!directory.exists()) {
```

```
        directory.mkdirs();
```

```
    }
```

```
    // Create a new writer for this batch
```

```
    ParquetWriter<GenericRecord> writer = AvroParquetWriter
```

```
        .<GenericRecord>builder(new Path(filePath))
```

```
        .withSchema(schema)
```

```
        .withCompressionCodec(CompressionCodecName.UNCOMPRESSED)
```

```
        .build();
```

```
    // Write all records in the batch
```

```
    long batchSize = 0;
```

```
    for (GenericRecord record : records) {
```

```
        writer.write(record);
```

```
        if (estimatedRecordSize == -1) {
```

```
            estimatedRecordSize = estimateRecordSize(record);
```

```
        }
```

```
        batchSize += estimatedRecordSize;
```

```
    }
```

```
    // Close the writer immediately after writing the batch
```

```
    writer.close();
```

```
    System.out.println("Wrote batch of " + records.size() + " records to " + filePath +  
        ", estimated size: " + (batchSize / (1024 * 1024)) + " MB");
```

```

}

private long estimateRecordSize(GenericRecord record) {
    // Estimate the size of the record by summing the sizes of its fields
    long size = 0;
    for (Schema.Field field : record.getSchema().getFields()) {
        Object value = record.get(field.name());
        if (value != null) {
            if (value instanceof CharSequence) {
                size += ((CharSequence) value).length() * 2; // Approximate size of a string
            } else if (value instanceof Integer) {
                size += Integer.BYTES;
            } else if (value instanceof Long) {
                size += Long.BYTES;
            } else if (value instanceof Float) {
                size += Float.BYTES;
            } else if (value instanceof Double) {
                size += Double.BYTES;
            } else if (value instanceof GenericRecord) {
                size += estimateRecordSize((GenericRecord) value); // Recursively estimate nested
records
            }
        }
    }
    return size;
}

private synchronized void closeAllWriters() {
    // Write any remaining batches when shutting down
    for (Map.Entry<String, List<GenericRecord>> entry : recordBatches.entrySet()) {
        if (!entry.getValue().isEmpty()) {
            try {
                writeToParquet(entry.getKey(), entry.getValue());
            } catch (IOException e) {
                System.err.println("Error writing final batch for " + entry.getKey() + ": " +
e.getMessage());
            }
        }
    }
    recordBatches.clear();
}

public void shutdown() {
    running.set(false);
}

```

```

/**
 * Checks for partitions that haven't received data for a while
 * and flushes any pending records to disk
 */
private void checkInactivePartitions() {
    long currentTime = System.currentTimeMillis();
    Set<String> partitionsToCheck = new HashSet<>(lastActivityTimestamp.keySet());

    for (String partitionKey : partitionsToCheck) {
        Long lastActivity = lastActivityTimestamp.get(partitionKey);
        if (lastActivity != null && (currentTime - lastActivity) > inactivityThresholdMs) {
            try {
                System.out.println("Partition " + partitionKey + " inactive for " +
                    (currentTime - lastActivity) / 1000 + " seconds. Flushing remaining records.");

                // Get buffered records for this partition
                List<GenericRecord> records = recordBatches.get(partitionKey);
                if (records != null && !records.isEmpty()) {
                    writeToParquet(partitionKey, records);
                    recordBatches.put(partitionKey, new ArrayList<>());
                }

                // Remove from activity tracking
                lastActivityTimestamp.remove(partitionKey);
            } catch (IOException e) {
                System.err.println("Error flushing inactive partition " + partitionKey + ": " +
                    e.getMessage());
            }
        }
    }
}

// Helper class to track writer and file size
private static class WriterInfo {
    private final ParquetWriter<GenericRecord> writer;
    private long currentSize;
    private final String filePath;

    public WriterInfo(ParquetWriter<GenericRecord> writer, long initialSize, String filePath) {
        this.writer = writer;
        this.currentSize = initialSize;
        this.filePath = filePath;
    }
}

```

```

public ParquetWriter<GenericRecord> getWriter() {
    return writer;
}

public long getCurrentSize() {
    return currentSize;
}

public void incrementSize(long bytes) {
    currentSize += bytes;
}

public String getFilePath() {
    return filePath;
}
}

// Example main method to demonstrate usage
public static void main(String[] args) {
    // Updated schema definition to match the new message format
    String schemaJson = "{\"type\":\"record\",\"name\":\"WeatherStationRecord\",\"fields\":["
        + "{\"name\":\"s_no\",\"type\":\"long\"},"
        + "{\"name\":\"battery_status\",\"type\":\"string\"},"
        + "{\"name\":\"status_timestamp\",\"type\":\"long\"},"
        + "{\"name\":\"weather\",\"type\":{\"type\":\"record\",\"name\":\"WeatherData\",\"fields\":["
        + "    {\"name\":\"humidity\",\"type\":\"int\"},"
        + "    {\"name\":\"temperature\",\"type\":\"int\"},"
        + "    {\"name\":\"wind_speed\",\"type\":\"int\"}"
        + "]}]}]";

    Schema schema = new Schema.Parser().parse(schemaJson);

    String bootstrapServer = System.getenv("BOOTSTRAP_SERVER");
    if (bootstrapServer == null) {
        System.err.println("BOOTSTRAP_SERVER environment variable is not set.");
        bootstrapServer = "localhost:9092"; // Default value
    }
    String topic = System.getenv("TOPIC");
    if (topic == null) {
        System.err.println("TOPIC environment variable is not set.");
        topic = "weather-station"; // Default value
    }
    String outputDir = System.getenv("OUTPUT_DIR");
    if (outputDir == null) {
        System.err.println("OUTPUT_DIR environment variable is not set.");
    }
}

```

```

    outputDir = "/out/weather-data"; // Default value
}
String maxFileSizeMB = System.getenv("MAX_FILE_SIZE_MB");
long maxFileSizeBytesMB = 1024; // Default to 1GB
if (maxFileSizeMB != null) {
    try {
        maxFileSizeBytesMB = Long.parseLong(maxFileSizeMB);
    } catch (NumberFormatException e) {
        System.err.println("Invalid MAX_FILE_SIZE_MB value, using default 1GB.");
    }
}
String batchSizeStr = System.getenv("BATCH_SIZE");
int batchSize = 1000; // Default to 10K records per batch
if (batchSizeStr != null) {
    try {
        batchSize = Integer.parseInt(batchSizeStr);
    } catch (NumberFormatException e) {
        System.err.println("Invalid BATCH_SIZE value, using default 10K.");
    }
}
}

```

```

KafkaToParquetWriter writer = new KafkaToParquetWriter(
    bootstrapServer,
    topic,
    "parquet-writer-group",
    outputDir,
    schema,
    maxFileSizeBytesMB,
    batchSize
);

// Add shutdown hook
Runtime.getRuntime().addShutdownHook(new Thread(writer::shutdown));

// Start processing
writer.start();
}
}

```

Docker file for each module:

1) Kafka-to-bitcask:

```
# Dockerfile for Normalizer
FROM eclipse-temurin:21-jdk

WORKDIR /app

COPY target/Normalizer-1.0-SNAPSHOT.jar /app/normalizer.jar

RUN mkdir -p /app/jfr-recordings

ENV KAFKA_BOOTSTRAP_SERVERS=kafka:9092
ENV KAFKA_TOPIC=weather-station
ENV BITCASK_SERVER_HOST=bitcask
ENV BITCASK_SERVER_PORT=5000

CMD ["java", "-XX:+FlightRecorder",
"-XX:StartFlightRecording=duration=1m,filename=/app/jfr-recordings/recording.jf
r,settings=profile", "-jar", "normalizer.jar"]
```

2) Rain-detector:

```
FROM eclipse-temurin:21-jdk

WORKDIR /app

# Copy the pre-built JAR file from your local machine to the container
COPY target/raining_detector-1.0-SNAPSHOT-jar-with-dependencies.jar .

ENV KAFKA_BOOTSTRAP_SERVERS=kafka:9092

# Create directory for JFR recordings
RUN mkdir -p /app/jfr-recordings

# Run the application with a one-minute JFR recording
CMD ["java", "-XX:+FlightRecorder",
"-XX:StartFlightRecording=duration=1m,filename=/app/jfr-recordings/recording.jf
r,settings=profile", "-jar", "raining_detector-1.0-SNAPSHOT-jar-with-dependencies
.jar"]
```


3) Parquet-maker:

```
FROM eclipse-temurin:21-jdk

WORKDIR /app

COPY target/parquet-maker-1.0-SNAPSHOT-jar-with-dependencies.jar
/app/parquet-maker.jar

ENV BOOTSTRAP_SERVER=kafka:9092
ENV TOPIC=weather-station
ENV OUTPUT_DIR=/app/output
ENV MAX_FILE_SIZE_MB=1024
ENV BATCH_SIZE=100

# Create directory for JFR recordings
RUN mkdir -p /app/jfr-recordings

# Enable Java Flight Recorder with continuous recording
# Start with JFR recording to file with 1-hour chunks, rotating up to 5 files
CMD ["java", \
    "-XX:+FlightRecorder", \
    "-XX:StartFlightRecording=duration=1m,filename=/app/jfr-recordings/recording.jf\
r,settings=profile", \
    "-jar", "parquet-maker.jar"]
```

Kubernetes yaml config files:

1) Bitcask deployment and service:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: bitcask
spec:
  replicas: 1
  selector:
    matchLabels:
      app: bitcask
  template:
    metadata:
      labels:
        app: bitcask
    spec:
      initContainers:
        - name: wait-for-pvc
          image: busybox:latest
          command: ['sh', '-c', 'until ls -la /data && touch /data/test-file && rm /data/test-file; do echo
"Waiting for PVC to be mounted..."; sleep 2; done; mkdir -p /data/logs /data/hints;echo "PVC is mounted
and writable!""']
          volumeMounts:
            - name: bitcask-data
              mountPath: /data
      containers:
        - name: bitcask
          image: bitcask:latest
          imagePullPolicy: Never
          ports:
            - containerPort: 5000
          volumeMounts:
            - name: bitcask-data
              mountPath: /data
      volumes:
        - name: bitcask-data
          persistentVolumeClaim:
            claimName: bitcask-pvc
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
```

```
name: bitcask-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
  storageClassName: standard
```

```
apiVersion: v1
kind: Service
metadata:
  name: bitcask
  labels:
    app: bitcask
spec:
  type: NodePort
  selector:
    app: bitcask
  ports:
    - protocol: TCP
      port: 5000
      targetPort: 5000
      name: bitcask-port
```

2) Kafka deployment and service:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kafka
spec:
  replicas: 1
  selector:
    matchLabels:
      app: kafka
  template:
    metadata:
      labels:
        app: kafka
    spec:
      containers:
        - name: kafka
          image: bitnami/kafka:latest
          resources:
            requests:
              memory: "1Gi"
              cpu: "500m"
            limits:
              memory: "2Gi"
              cpu: "1000m"
      env:
        - name: KAFKA_CFG_AUTO_CREATE_TOPICS_ENABLE
          value: "true"
        - name: KAFKA_CFG_LISTENER_SECURITY_PROTOCOL_MAP
          value: "INTERNAL:PLAINTEXT,EXTERNAL:PLAINTEXT,CONTROLLER:PLAINTEXT"
        - name: KAFKA_CFG_CONTROLLER_LISTENER_NAMES
          value: "CONTROLLER"
        - name: KAFKA_KRAFT_CLUSTER_ID
          value: "kimokono"
        - name: ALLOW_PLAINTEXT_LISTENER
          value: "yes"
        - name: KAFKA_CFG_LISTENERS
          value: "INTERNAL://:9092,EXTERNAL://:29092,CONTROLLER://:9093"
        - name: KAFKA_CFG_CONTROLLER_QUORUM_VOTERS
          value: "0@kafka:9093"
        - name: KAFKA_CFG_ADVERTISED_LISTENERS
          value: "INTERNAL://kafka:9092,EXTERNAL://localhost:29092"
        - name: KAFKA_ENABLE_KRAFT
          value: "yes"
```

```
- name: KAFKA_CFG_PROCESS_ROLES
  value: "broker,controller"
- name: KAFKA_CFG_INTER_BROKER_LISTENER_NAME
  value: "INTERNAL"
- name: KAFKA_CFG_NODE_ID
  value: "0"
ports:
- containerPort: 9092
- containerPort: 9093
- containerPort: 29092
volumeMounts:
- name: kafka-data
  mountPath: /bitnami/kafka
volumes:
- name: kafka-data
  persistentVolumeClaim:
    claimName: kafka-pvc
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: kafka-pvc
spec:
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
  storageClassName: standard
```

```
apiVersion: v1
kind: Service
metadata:
  name: kafka
spec:
  type: NodePort
  ports:
  - port: 9092
    targetPort: 9092
    name: internal
  - port: 9093
    targetPort: 9093
    name: controller
  - port: 29092
    targetPort: 29092
```

```
name: external
selector:
  app: kafka
```

3) bitcask message normalizer deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: bitcask-normalizer
spec:
  selector:
    matchLabels:
      app: bitcask-normalizer
  template:
    metadata:
      labels:
        app: bitcask-normalizer
    spec:
      containers:
        - name: bitcask-normalizer
          image: bitcask-normalizer:latest
          imagePullPolicy: Never
          env:
            - name: KAFKA_BOOTSTRAP_SERVERS
              value: "kafka:9092"
            - name: KAFKA_TOPIC
              value: "weather-station"
            - name: BITCASK_SERVER_HOST
              value: "bitcask"
            - name: BITCASK_SERVER_PORT
              value: "5000"
```

4) parquet-maker deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: parquet-maker
spec:
  selector:
    matchLabels:
      app: parquet-maker
  template:
    metadata:
      labels:
        app: parquet-maker
    spec:
      initContainers:
        - name: init-parquet-maker
          image: busybox
          command: ['sh', '-c', 'until ls -la /app/output; do echo "Waiting for persistent storage to be ready..."; sleep 2; done;']
          volumeMounts:
            - name: parquet-data
              mountPath: /app/output
      containers:
        - name: parquet-maker
          image: parquet-maker:latest
          imagePullPolicy: Never
          env:
            - name: BOOTSTRAP_SERVER
              value: "kafka:9092"
            - name: TOPIC
              value: "weather-station"
            - name: OUTPUT_DIR
              value: "/app/output"
            - name: MAX_FILE_SIZE_MB
              value: "1"
            - name: BATCH_SIZE
              value: "1000"
          volumeMounts:
            - name: parquet-data
              mountPath: /app/output
      volumes:
        - name: parquet-data
          persistentVolumeClaim:
            claimName: parquet-data-pvc
```

```
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: parquet-data-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
  storageClassName: standard
---
```

5) Rain detector deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: rain-detector
spec:
  replicas: 1
  selector:
    matchLabels:
      app: rain-detector
  template:
    metadata:
      labels:
        app: rain-detector
    spec:
      containers:
        - name: rain-detector
          image: rain-detector:latest
          imagePullPolicy: Never
          env:
            - name: KAFKA_BOOTSTRAP_SERVERS
              value: "kafka:9092"
```


6) Weather-station statful-set:

Used statefulSet mainly for the naming of the pods only as weather stations don't need to be stateful or keep any data

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: weather-station
spec:
  serviceName: "weather-station"
  replicas: 10
  selector:
    matchLabels:
      app: weather-station
  template:
    metadata:
      labels:
        app: weather-station
    spec:
      containers:
        - name: weather-station
          image: weather-station:latest
          imagePullPolicy: Never
          env:
            - name: KAFKA_BOOTSTRAP_SERVERS
              value: "kafka:9092"
            - name: STATION_ID
              valueFrom:
                fieldRef:
                  fieldPath: metadata.name
---
apiVersion: v1
kind: Service
metadata:
  name: weather-station
  labels:
    app: weather-station
spec:
  clusterIP: None
  selector:
    app: weather-station
```

7) elasticsearch and kibana deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: elasticsearch
spec:
  replicas: 1
  selector:
    matchLabels:
      app: elasticsearch
  template:
    metadata:
      labels:
        app: elasticsearch
    spec:
      initContainers:
        - name: init-sysctl
          image: busybox:latest
          command:
            - sysctl
            - -w
            - vm.max_map_count=262144
          securityContext:
            privileged: true
      containers:
        - name: elasticsearch
          image: nshou/elasticsearch-kibana
          resources:
            requests:
              memory: "1Gi"
              cpu: "500m"
            limits:
              memory: "2Gi"
              cpu: "1000m"
          ports:
            - containerPort: 9200
            - containerPort: 5601
          env:
            - name: SSL_MODE
              value: "false"
            - name: ES_JAVA_OPTS
              value: "-Xms512m -Xmx512m"
            - name: discovery.type
              value: "single-node"
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: elasticsearch
spec:
  type: NodePort
  ports:
    - port: 9200
      targetPort: 9200
      name: elasticsearch
    - port: 5601
      targetPort: 5601
      name: kibana
  selector:
    app: elasticsearch
```

8) parquet-to-es cronJob:

Used to load newly added file to the persistent-volume of the parquet maker pod into elastic search after transforming data for the desired business needs.

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: parquet-to-elastic
spec:
  schedule: "*/30 * * * *" # Runs every half an hour
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: parquet-to-elastic
              image: parquet-to-es:latest
              imagePullPolicy: Never
              env:
                - name: PARQUET_DIR
                  value: "/app/weather-data"
                - name: PROCESSED_FILES
                  value: "/app/out/processed_files.txt"
                - name: ELASTIC_HOST
                  value: "elasticsearch" # Elasticsearch service name
                - name: ELASTIC_PORT
                  value: "9200"
                - name: ELASTIC_INDEX
                  value: "weather_stations_metrics"
              volumeMounts:
                - name: parquet-data
                  mountPath: "/app/weather-data"
                - name: processed-files
                  mountPath: "/app/out"
            restartPolicy: OnFailure
          volumes:
            - name: parquet-data
              persistentVolumeClaim:
                claimName: parquet-data-pvc
            - name: processed-files
              persistentVolumeClaim:
                claimName: processed-files-pvc
  ---
apiVersion: v1
```

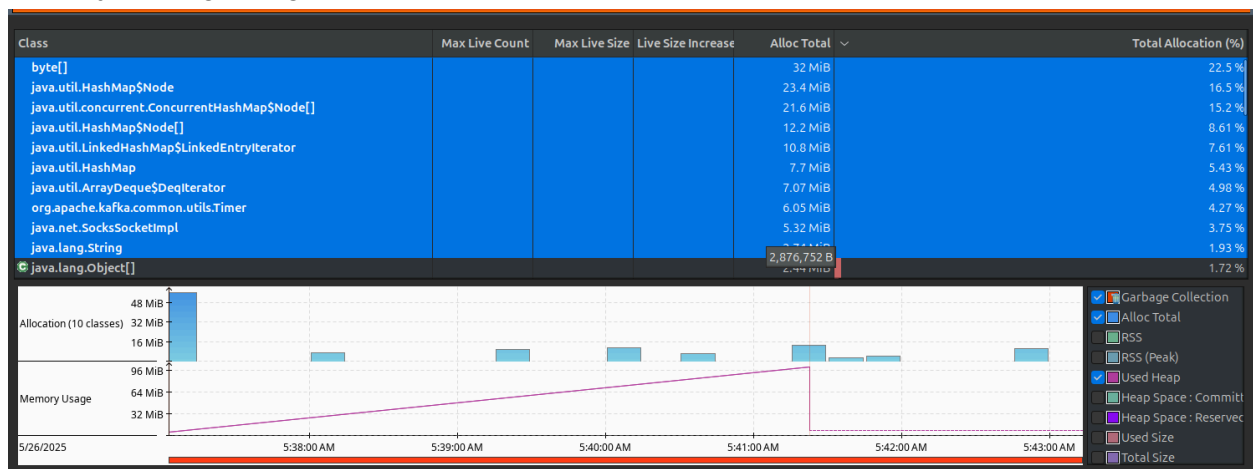
```
kind: PersistentVolumeClaim
metadata:
  name: processed-files-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi # Adjust size as needed
  storageClassName: standard
---
```

Profiling center station using JFR(java flight recorder)

The **Central station** consists of 3 main but separate modules (message to bitcask normalizer + parquet-maker + rain-detector module) each module does one job and has its own consumer group in kafka (they are connected to weather stations in pub-sub pattern)
Run was for nearly 3 minutes to ensure IO writes.

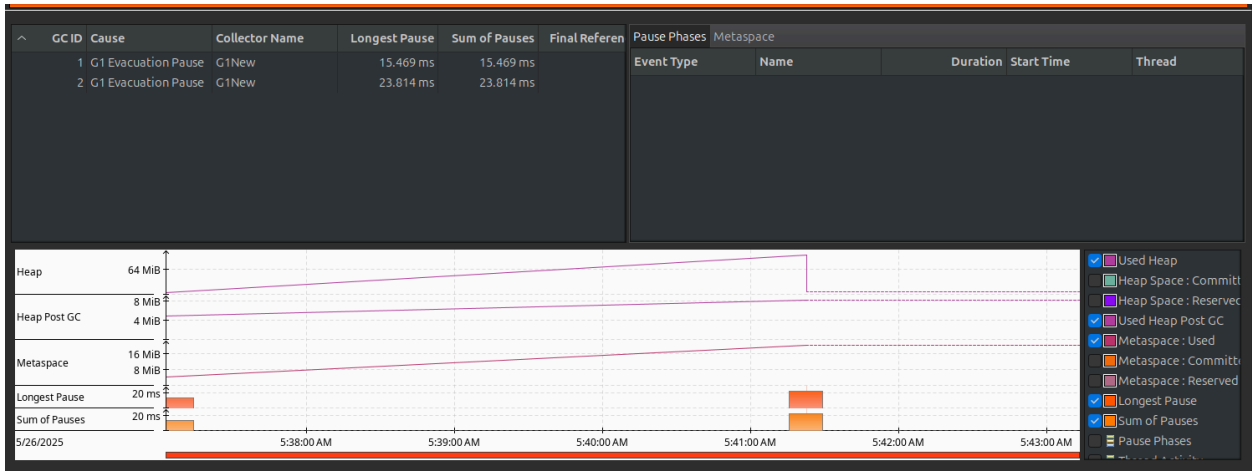
Main points :

- Top 10 Classes with highest total memory
 1. byte[]
 2. java.util.HashMap\$Node
 3. java.util.concurrent.ConcurrentHashMap\$Node[]
 4. java.util.HashMap\$Node[]
 5. java.util.LinkedHashMap\$LinkedEntryIterator
 6. java.util.HashMap
 7. java.util.ArrayDeque\$DeqIterator
 8. org.apache.kafka.common.utils.Timer
 9. java.net.SocksSocketImpl
 10. java.lang.String



- GC pauses count

2 pause counts

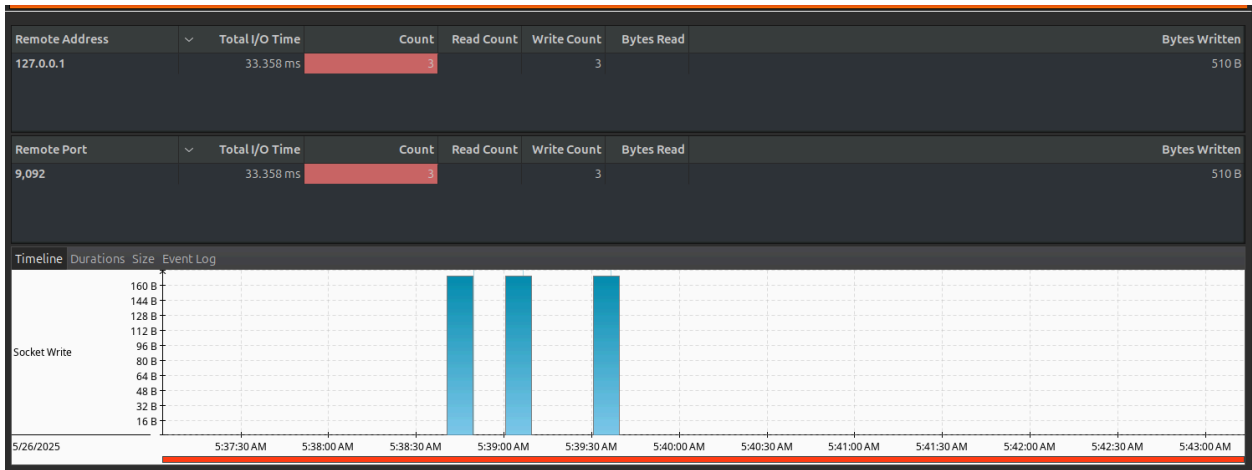


- GC maximum pause duration

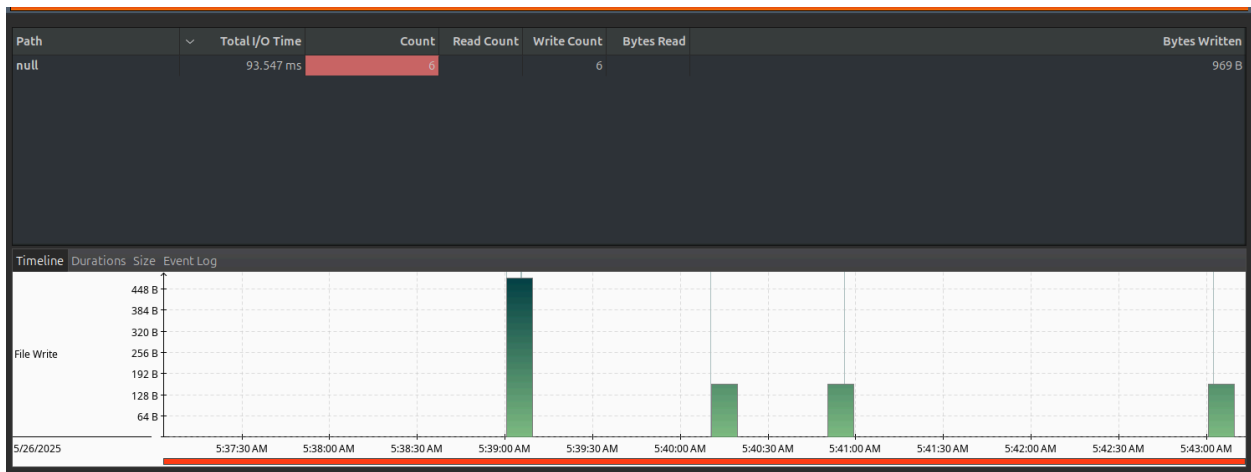
23.814 ms

- List of I/O operations

3 socket I/O operations each one takes the socket time of 10 ms sending messages



6 file I/O operations:



Bonus:

Integration patterns used in this project:

1) Channel adapter:

A message normalizer module in the central station provides the capability for the bitcask server to poll message from the kafka topic and save the latest value in its file system

2) Invalid message channel:

For modules of two module of the central station if a message has a missing format or a wrong json was passed inside the weather-station topic the module invalidates that message and passes it to another topic for invalid messages

3) Content enricher:

Using a separate module that runs as a cronJob on k8s we modify the parquet files produced by parquet-maker to add some points for analytics (show the percentage of dropped messages)

4) Message Normalizer:

If we can consider a message normalizer as the basic form of message translator then the message-to-bitcask module can be considered a message normalizer, since it can be extended to allow for other weather stations to send messages in different (should be supported) schema in the same channel and it will be kept in bitcask

5) Content filter:

The rain-detector module sends a more concise message through its channel where more modules can connect to it and use the filtered message

Example rain-detector message:

Rain detected at station ID: 1 at humidity level: 76
--

6) Polling consumer:

The kafka consumer api acts as a polling consumer where it keeps receiving messages and answers to poll() call from receiver

7) Durable subscriber:

We can argue that the the elastic-search pod may act as a subscriber to message summary that are kept in the form of .parquet files and more message will be produced even if the “subscriber” is not up